

Relational Database Model

Relational Model

- One single data modeling “tool”: relation, or a 2D table;
- A relational database is a collection of relations;
- High degree of data independence
- Association between information elements (constraints)

More Formally

Relation: a two-dimensional table of columns and rows.

Attribute, Field: name of a *column* in the relation.

- take values from predefined *domains*

Record, tuple: a single row in the relation: a collection of *attribute values*.

Schema: the name of a *relation* plus the set of *attributes* of the relation (and their domains).

- E.g. Book(ISBN string, Title string, Author string, year integer).

Relation instance: a set of *tuples* for a given *relation*.

- changes with time (as stuff gets added, deleted, modified)
- schema usually does not change (although it might in some cases)

Cardinality: number of tuples in a relation

Degree: number of attributes in a relation

Constraints

Superkey a collection of attributes in a relation that *uniquely identifies each tuple in it*.

Candidate key a *superkey* that has no *superkey subsets*

Primary key one *candidate key* per relation, designated to be the main way of maintaining tuple uniqueness.

Key constraint : each relation must have a primary key.

Foreign key a *primary key* of one relation, included in the attributes of another relation (usually for the purpose of linking two components of the database together).

Referential integrity constraint each collection of values of a *foreign key* in a relation must appear as a *primary key* in the referenced relation.

Null value : a “no value” value for a relational attribute. Lack of value, or value not yet available.

not null constraint : a statement that a specific attribute is not allowed to have null values. (e.g., primary key attributes).

Examples

Example 1. A university registrar wants to organize a list of courses offered at a university. The information available about each course is its prefix (e.g., 'CSC'), its number (e.g., 365), its title (e.g., 'Introduction to Databases'), its catalog description, and the department responsible for scheduling the course. The registrar creates the following relational table to represent the courses:

Courses(Prefix String, CourseNumber Integer, Title String,
Catalog String, Department String)

A portion of this table would look as follows:

Prefix	CourseNumber	Title	Catalog	Department
CSC	365	Introduction to Databases	Basic principles of database management...	CSSE
CSC	357	Systems Programming	C programming language from a system...	CSSE
CHEM	212	Introduction to Organic Chemistry	Structure, isomerism, nomenclature, fundamental...	Chemistry
STAT	331	Statistical Computing with R	Importing, managing, and cleaning data from...	Statistics
STAT	365	Statistical Communication	Written communication of statistical ideas...	Statistics

Determining the keys:

- The pair of attributes (Prefix, CourseNumber) *together* is a candidate key for the Courses table. The requirement that the Prefix-CourseNumber pair is unique for each course comes from *outside* - this is university catalog rule established before the Courses table is introduced.
- The Prefix attribute alone is not enough to uniquely identify every row in the table (e.g., there are two that have the CSC prefix).

- The **CourseNumber** attribute alone is also not enough: there are two rows describing different courses whose number is 365.
- The second candidate key in the table could be **Title**. This means that the Registrar's office requires all courses to have a unique title.
- Between the **Title** and the (**Prefix**, **CourseNumber**), the registrar selects (**Prefix**, **CourseNumber**) as the **primary key**.

Example 2. After examining the **Courses** table, the registrar arrives to the conclusion that they want another attribute in the table that merges (concatenates) the course prefix and the course number¹. As a result, the new **Courses** table looks as follows:

```
Courses(CourseId String, Prefix String, CourseNo Integer,
        Title String, Catalog String, Department String)
```

A portion of this table looks as follows:

CourseId	Prefix	CourseNo	Title	Catalog	Department
CSC 365	CSC	365	Introduction to Databases	Basic principles of database management. . .	CSSE
CSC 357	CSC	357	Systems Programming	C programming language from a system. . .	CSSE
CHEM 212	CHEM	212	Introduction to Organic Chemistry	Structure, isomerism, nomenclature. . .	Chemistry
STAT 331	STAT	331	Statistical Computing with R	Importing, managing, and cleaning data . . .	Statistics
STAT 365	STAT	365	Statistical Communication	Written communication of statistical ideas. . .	Statistics

We now have the following candidate keys in this table:

- (**Prefix**, **CourseNo**) - our old primary key
- **Title** - per our assumption that the Registrar's office will not let two courses with exactly the same name to co-exist.
- **CourseId** - as this is literally concatenation of the two values that forms a candidate key.

Given these three candidate keys, the registrar now selects **CourseId** as the **primary key** for this table.

Example 3. In addition to the courses, the registrar wants to organize a list of the departments. The registrar needs the following information about each department: its full name, the college it is in, its location on campus (building number and room number). The first version of the **Departments** table looks as follows:

```
Departments(Name String, College String, Building Integer, Room String)
```

(note: rooms may have values like '245a' so, they are strings.)

A portion of the table looks as follows:

Name	College	Building	Room
Computer Science and Software Engineering	CENG	14	245
Statistics	COSAM	25	107D
Chemistry and Biochemistry	COSAM	180	206
Physics	COSAM	180	207

Keys:

¹This is because there is a lot of code being written that accesses the table and has to concatenate these two attributes every single time.

- The **Name** attribute is a candidate key (no two departments can have the same name).
- Assuming no two departments share the same office - a reasonable assumption at Cal Poly, the pair (**Building**, **Room**) is also a candidate key.
- These are the only candidate keys in the table (**College** by itself is NOT a candidate key, and any other combination of attributes either contains a key, or is not unique - e.g., (**College**, **Building**) is not a candidate key, because multiple departments can have offices in the same building, and (**College**, **Room**) is not unique because multiple departments can have offices with the same room number but in different buildings).
- Between the two candidate keys, the registrar selects **Name** as the **primary key**.